

Probabilistic Models of time series data

Dirichlet Processes

- Krishna Devkota

Motivation

- Your friends are always complaining about food in the mensa
- So, you decide to start a restaurant nearby
- But, there is a problem ...
 - you have no clue about the food preferences of the students (*some can't do without beer and pizza, some prefer vegan tofu, some want sushi maybe*)
- So, you start with a survey of people around you
- **Goal:**
 - Find the food preferences of the students and cluster them into appropriate groups

Clustering problem:

- Given observations: x_1, x_2, \dots, x_n

Objective:

- subdivide them into subsets, i.e. **clusters**
- some sort of **similarity** sought for observations within each **cluster**




Fig: 1000 points divided into 3 clusters

Clustering in terms of mixture models:


First, some basics!

Modeling assumption:

- Data is partitioned into groups / clusters
- Each observation x_i belongs to a single cluster k

Express **cluster assignment** as **random variable** L_i

- so, $L_i = k$ means X_i belongs to cluster k
- As cluster assignments unknown, L_i remains *latent* (*unobserved*)




Stating it differently,

what is *unknown* is the **partition** of data-space $\{1, \dots, N\}$

N = number of observed data-points

- Distribution characterizing a **single cluster** k , given as:

$$P_k(\bullet) := \mathbb{P}[X \in \bullet | L = k]$$



- Probability that newly generated observation belongs to cluster k (*basically the weights*):

$$c_k := \mathbb{P}\{L = k\}$$

As \mathbf{c}_k are probabilities of *mutually exclusive events*, they sum up to 1.

So, now the distribution of \mathbf{X} is of the form:

$$P(\bullet) = \sum_{k \in \mathbb{N}} c_k P_k(\bullet)$$

- A *model* of this form is called a **mixture distribution**.




Figure: mixture of two Gaussian, left: $c_1 = 0.5$, $c_2 = 0.5$, right: $c_1 = 0.8$, $c_2 = 0.2$

- If the **number** of **clusters** is finite, meaning if there are finite number of K with **non-zero probabilities** c_k , we call the mixture a **finite mixture**

Mixture models:

$$p(x|m) = \int_{\Omega_{\theta}} p(x|\theta)m(d\theta)$$

Where, $p(x|\theta)$ is the component distribution

$m(d\theta)$ is called the *mixing measure*

To see what this means,

Think of $p(x|\theta)$ as a *Gaussian distribution* with parameter θ ,

- The second part $m(d\theta)$ is a distribution over θ
- $m(d\theta)$ is a distribution of finitely many *delta spikes*.

($p(x|\theta)$ defines what kind of mixture model it is, i.e. if $p(x|\theta)$ is a Gaussian then a Gaussian mixture model)

Mixture models can be looked upon as a two stage sampling:

For a given mixture model:

$$p(x|m) = \int_{\Omega_\theta} p(x|\theta)m(d\theta)$$

- first sample θ from the *mixing measure* ' m '
- plug-in the theta as a parameter of the distribution $p(\mathbf{x}|\theta)$
- then sample the random variable \mathbf{x} from the resulting distribution $p(\mathbf{x}|\theta)$

Sample $X \sim p(\cdot | m)$ as:

1. $\Theta \sim m$
2. $X \sim p(\cdot | \theta)$

Parameter space for a mixture model:

- Set of all *discrete* probabilities on the *parameter space* Ω_ϕ defined by $p(\mathbf{x}|\phi)$
- Φ = location of *atoms*

Finite mixture model

$$p(x|\boldsymbol{\theta}, \mathbf{c}) = \int_{\Omega_{\theta}} p(x|\theta) m(d\theta) \quad \text{with} \quad m(\cdot) = \sum_{k=1}^K c_k \delta_{\theta_k}(\cdot)$$

- where, mixing measure $m(\cdot)$ is simply **sum** over finite number of delta **spikes** at k different locations

- \mathbf{c}_k 's are the *convex coefficients* to weigh the **Dirac deltas**
(convex coefficients simply mean that c values are non-negative and sum to 1)

Note: The mixture model is parameterized by the mixing distribution/ mixing measure!

Bayesian mixture model

A **Bayesian mixture model** is simply a **mixture model** with a **random mixing measure**.

Randomizing our *parameters* \mathbf{c} and $\boldsymbol{\theta}$ will give us a random *mixing measure*:

$$M(\cdot) = \sum_{k=1}^K C_k \delta_{\Theta_k}(\cdot)$$

We basically want to put a prior for the mixing distribution!

How to choose the priors?

Idea from the very start: Use a **conjugate prior**!

$$M(\cdot) = \sum_{k=1}^K C_k \delta_{\Theta_k}(\cdot)$$

What is a **conjugate prior**?

- **Posterior** belongs to the same class of distribution as the *prior*!
- **Conjugate priors** generally occur only in the **exponential families** in the **parametric** case!

This accounts for the θ 's but we still have the C 's!

$$M(.) = \sum_{k=1}^K C_k \delta_{\Theta_k}(.)$$

What could be the distribution of c's?

Observation:

- when sampling from the ***mixture-model*** the c's (*index of the clusters*) are found to be: ***multinomial distributed***
- The parameter of this distribution is the *vector of coefficients* [c1,c2,...ck]

The conjugate prior of a multinomial is a Dirichlet!

So, we'll use a Dirichlet distribution on our weight vector [c1,c2,...ck]

Let's quickly look back at our restaurant problem:

We could use one of the standard algorithms based on *FMM*:

- k-means clustering (*not a probabilistic model*)
 - Gaussian mixture models
-
- But remember, we surveyed very limited group of students!
 - In reality, there are infinite (well, very large) number of foodies out there!
 - How do you account for *new food choices* of other students (***hidden clusters***) which may arise later on?

To use a *Finite mixture model*, we (need to) start by *assuming* a certain (finite) number of clusters

But, what we really want is to *allow newer clusters* as our data (*students*) grow!

So, we have to do it in a non-parametric way!

Very quick look at **parametric** vs. **non-parametric** model

Parametric:


Number of parameters fixed w.r.t sample size

Non-parametric:


- number of parameters grow with sample size
- infinite dimensional parameter space

(parameter space: space of models that can explain our data)

Quick example of density estimation:



Parametric



Nonparametric

Note: Fig. not necessarily a Bayesian!

Dirichlet process mixture: Infinite limit of finite mixture models:

In the *Bayesian mixture model*, the *mixing measure* was:

$$M(\cdot) = \sum_{k=1}^K C_k \delta_{\Theta_k}(\cdot)$$

If $K \rightarrow \infty$, in our previous model, we get a **non parametric** model :

$$M(\cdot) = \sum_{k=1}^{\infty} C_k \delta_{\Theta_k}(\cdot) \quad \text{where} \quad \sum_{k=1}^{\infty} C_k = 1$$

When nonparametric, we need to generate $M(\cdot)$ at random. Think about how we could do that?

Also, note the restriction: *infinite sum* over \mathbf{c} 's must be 1.

Simple distribution that does this: Dirichlet process!

What do we mean by simple?

- as much *independence* as possible

(if c 's and θ 's *highly dependent*, very difficult to do inference on that model)

Dirichlet Process:

A Dirichlet process is a distribution on random probability measures of the form:

$$M(\cdot) = \sum_{k=1}^{\infty} C_k \delta_{\Theta_k}(\cdot) \quad \text{where} \quad \sum_{k=1}^{\infty} C_k = 1$$

Take some base distribution \mathbf{G}_0 and sample *atoms* θ 's iid from it (*infinite seq of thetas*)

What about the weights?

- they can't be independent, because should sum up to one!
(so if 1st weight 0.5, second can't be greater than 0.5)

Next best thing we can do?


- sample *independent proportions* and use them to generate weights.

This brings us to: *stick breaking construction!*

Constructive definition of $DP(\alpha, G_\varphi)$ [stick-breaking analogy]

STEPS:

- Start with a stick of length one
- Generate a random variable: $\beta_1 \sim \text{Beta}(1, \alpha)$
- By the definition of the Beta distribution, this will be a real number between 0 and 1, with expected value $1/(1+\alpha)$
- Break off the stick at b_1
- $w_1 = 1 - b_1$ is then the length of the stick on the left
- Now, generate $\beta_2 \sim \text{Beta}(1, \alpha)$
- Break off the part of stick b_2
- Again, w_2 is the length of the stick to the left, i.e., $w_2 = (1 - b_1)b_2$
- Repeat the process.



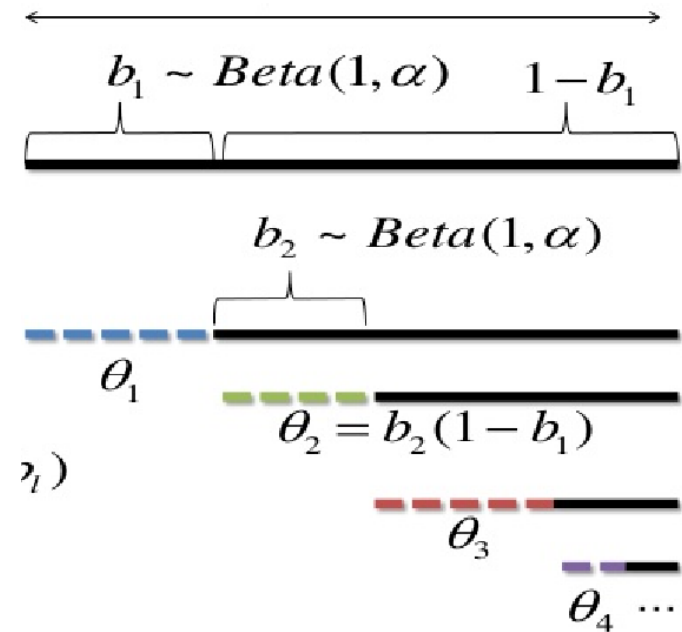
Mathematically,
the process can be described as:

$$\Theta_k \sim_{\text{iid}} G_0$$

$$V_k \sim_{\text{iid}} \text{Beta}(1, \alpha)$$

Compute C_k as

$$C_k := V_k \prod_{i=1}^{k-1} (1 - V_i)$$




Dirichlet process in our clustering problem:


$$M(\cdot) = \sum_{k=1}^{\infty} C_k \delta_{\Theta_k}(\cdot) \quad \text{where} \quad \sum_{k=1}^{\infty} C_k = 1$$

Intuitive understanding of the above expression:

Let's say the centers of our clusters come from a base distribution (Go) that is a Gaussian:



- each line here is a weighted atom
- each draw here will be our cluster
- these are the center of our clusters



- For each atoms in the left figure, new Gaussian is created around the point that we drew

If our data comes from **cluster 1**, lands in **first area** with **high probability**, and so on!

So far we,

- described a *generative model* that allows us to calculate probability of assigning any particular *set of groups* to our *data points*

But, how do we actually learn a good *set of group assignments*!

In other words, how do we infer the model!

Inference in the DP mixture (Gibbs sampling):

Approach: MCMC sampling (particularly Gibbs)


Three unknowns:

π (parameter), z_i (latent variable), θ_k^* (parameter)

Observation:

x_i (data)

Objective: learn about π , z_i , θ_k^* in an iterative fashion



Inference (Gibbs sampling):

STEPS:

Step 1:

- Start with some estimate of parameters θ_k^* and π
- Compute the “conditional distribution” of z_i given the parameters and data
- Draw a sample from that conditional distribution
(*hope: to achieve high probability values of z_i that are more likely given the param and data*)

Step 2:


- After sampling the latent variables (z_i), conditioned on these latent variables and our data, compute the conditional distribution of the parameters

Step 3:

- Sample from the newly computed distribution of the parameters, and the update the value of latent variables.
- This completes one cycle of *MCMC* update.

Repeat the above steps till convergence!

Will converge to true posterior distribution!




Chinese Restaurant Process:

The distribution of **random partition** induced by the [Dirichlet process](#) can be clearly explained through the [Chinese Restaurant Process](#).

Quick Definition of Partition:


A partition of a set S is defined as a disjoint family of non-empty subsets of S whose union is S .

- $S = \{\text{Alice, Bob, Charles, David, Emma, Florence}\}.$
- $\varrho = \{ \{\text{Alice, David}\}, \{\text{Bob, Charles, Emma}\}, \{\text{Florence}\} \}.$



- Denote the set of all partitions of S as \mathcal{P}_S .
- **Random partitions** are random variables taking values in \mathcal{P}_S .

Chinese Restaurant Process



Each customer comes into a restaurant and sits at a table

$$\mathbb{P}(\text{sit at table } c) = \frac{n_c}{\alpha + \sum_{c \in \varrho} n_c}$$


$$\mathbb{P}(\text{sit at new table}) = \frac{\alpha}{\alpha + \sum_{c \in \varrho} n_c}$$

Customers correspond to elements of : S

Tables correspond to clusters in : ϱ

Rich get richer : large clusters more likely to attract new customers

Number of clusters of a DP:



If K_n is number of clusters in sample of size n , then:

$$\mathbb{E}[K_n] = O(\log(n))$$

Chinese Restaurant Process

Exchangeable Partition Probability function:

Multiplying conditional probabilities together, the overall probability of \mathcal{Q}

$$\mathbb{P}(\mathcal{Q}|\alpha) = \frac{\alpha^{|\mathcal{Q}|} \Gamma(\alpha)}{\Gamma(n + \alpha)} \prod_{c \in \mathcal{Q}} \Gamma(|c|)$$

This is known as Exchangeable Partition Probability function.

We see that the ***order of the customers entering*** the restaurant ***does not affect*** the probability of \mathcal{Q}

As we see above, ex-changeability is one of the important properties of Dirichlet process!!

What we have seen so far:

Non-parametric Bayesian clustering:

- Infinite number of clusters $K_n \leq n$, where n = observations
- If partition exchangeable, it can be modeled using a random discrete distribution

Inference:

- since partitions (cluster assignments) not directly observed, latent variable algorithm used
- Gibbs sampling

Assumptions:

- prior assumption of number of clusters K_n
- distribution of cluster sizes

Dynamic Hierarchical Dirichlet Process (Extend DP to incorporate time dependence)

But first,

Hierarchical Dirichlet Process:

Dirichlet Process where the base distribution \mathbf{G}_0 is itself drawn from the *Dirichlet Process*.

This allows to *share a common parameter* across the ***groups*** through the base distribution, even if the different groups have their own distribution.

To apply the Dirichlet process to a time-evolving sequence of data

Assumption:

(i) Two data samples drawn at proximate times

- have a higher probability of sharing the same underlying model parameters (atoms) than parameters drawn at disparate times

(ii) Possibility that temporally distant data samples may also share model parameters,

- Possible distant repetition in the data

Statistical properties of data collected at consecutive time points are linked

BUT how?

Steps:

- Use a *random parameter* that controls their probabilistic similarity.
- Derive a sharing mechanism of time evolving data
- then, develop an appropriate Markov Chain Monte Carlo (MCMC) sampler

Possible applications:

- Music segmentation
- Gene expression data

Experimental Result:

Music Segmentation:

In music segmentation:

- interesting to infer relationship between different parts of a piece
- also similarity between different pieces

So,

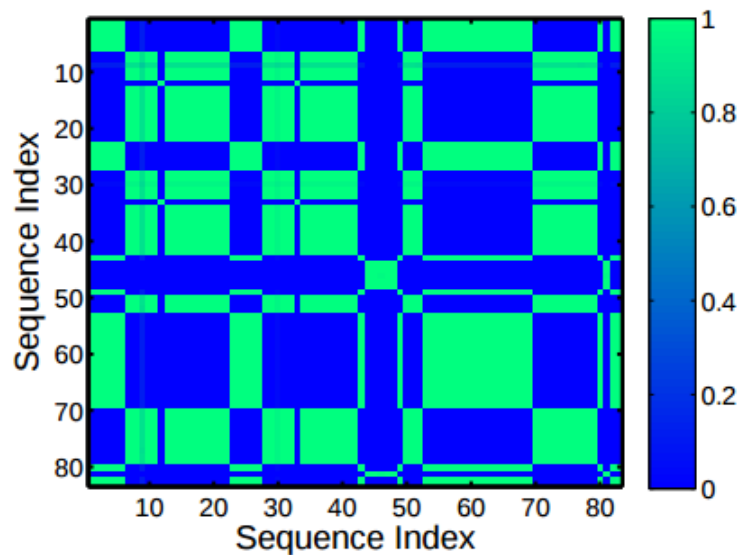
- A music piece *divided* into different *contiguous sub-sequences*
- Each *subsequence* modeled via a **HMM**
- **dHDP** useful here in enforcing the idea that *contiguous sub-sequences* are likely to *fall within the same music segment*, and thus *share HMM parameters*
- when the *segment changes*, changes *detected* by *dHDP*

STEPS:


- First, ***MFCC features*** extracted
- ***discretized*** with ***vector quantization***
- Piece transformed into ***4980 discrete symbols***, and ***83 subsequences***
- Each ***subsequence*** is modeled using ***HMM*** with ***8 states***
- Each subsequence ***6 sec*** in length

To model the time dependence between adjacent subsequences,

- each ***subsequence*** corresponds to ***one group*** in the ***dHDP HMM*** mixture
- choose ***one set of HMM parameters*** according to the ***corresponding mixture weights***
- In the ***dHDP framework***, one subsequence can ***share*** the ***old DP mixture distributions*** with the ***previous ones***
- Or it might be ***drawn*** from an ***innovation DP mixture***,
- this may be also ***shared*** by the ***following time series*** in a similar manner



(a)



(b)

Figure 3. Similarity matrix $E(z'z)$ from HMM mixture modeling of the Sonata. (a) dHDP-HMM, (b) HDP-HMMs.

Conclusion:

- Clustering problem
- Mixture models

Finite mixture model

Infinite mixture model

-Dirichlet as a prior for IMM

- Dirichlet Process
- Construction of Dirichlet mixtures (*using stick breaking analogy*)
- Dirichlet mixture Inference
MCMC sampling (Gibbs sampling)
- Random partitions introduced by Dirichlet process
(Explained through Chinese Restaurant process)

- Properties of Dirichlet Process
 - Exchangeability
- Hierarchical Dirichlet Process
- Hierarchical Dirichlet Process for Dynamic sequence
 - Dynamic Hierarchical Dirichlet Process*
(sharing some parameters across the groups)
- Application of dHDP
 - Music segmentation*

References:

Lu Ren, David B. Dunson, and Lawrence Carin. 2008. The dynamic hierarchical Dirichlet process. In *Proceedings of the 25th international conference on Machine learning (ICML '08)*. ACM, New York, NY, USA, 824-831. DOI=<http://dx.doi.org/10.1145/1390156.1390260>

Lecture notes on Bayesian Nonparametrics, Peter Orbanz
http://stat.columbia.edu/~porbanz/papers/porbanz_BNP_draft.pdf

http://stat.columbia.edu/~porbanz/talks/MLSS12_1.pdf

Bayesian Nonparametrics, Yee Whye Teh, Dept of Statistics, Oxford
http://mlss.tuebingen.mpg.de/2013/2013/slides_teh.pdf

**Questions /
Comments ??**